



16×4 SRAM Memory Design

22 nm High-Performance CMOS

ESE 3700

Circuit Design and Optimization

Tyrone Marhguy

Source repository: github.com/tmarhguy/64b-sram

Academic Integrity Statement

I, **Tyrone Marhguy**, certify that I have complied with the University of Pennsylvania's Code of Academic Integrity in completing this project.

Contents

1	Executive Summary	1
1.1	Required vs. Delivered (headline)	1
1.2	Map of the report	2
2	Architecture and FOM Strategy	3
2.1	Reading the FOM	3
2.2	Top-level array organization: 16×4, no column mux	4
2.3	Subsystem-level decisions	5
2.4	Block diagram of the full memory	5
3	Primitives: building blocks the rest of the design instantiates	7
3.1	Inverter (and its skewed variants)	8
3.2	NAND2 and NAND3	8
3.3	Transmission gate	9
3.4	D flip-flop (transmission-gate master-slave)	9
4	Bitcell Design	11
4.1	Topology: 6T	11
4.2	Sizing and rationale	12
4.2.1	Cell ratio: read static-noise margin	13
4.2.2	Pull-up ratio: write margin	13
4.2.3	Why not minimum everything?	13
4.3	Bitline capacitance	14
4.4	Cell-level validation simulations: Tests 1–5	15
4.4.1	Test 1 — Write 1→0	15
4.4.2	Test 2 — Write 0→1	15
4.4.3	Test 3 — Read stability storing 0	16
4.4.4	Test 4 — Read stability storing 1	16
4.4.5	Test 5 — Data retention (hold)	17
5	Read Path (Column slice, read direction)	18
5.1	Precharge and equalize	19
5.2	Bitcell BL discharge	20
5.3	StrongARM sense amplifier	20
5.4	Replica column for self-timed SAE	22
5.4.1	Motivation	22
5.4.2	Replica circuit	23
5.4.3	Robustness	25
5.4.4	Bitcell-area accounting note	25
5.5	Output latch	25
6	Write Path (Column slice, write direction)	26
6.1	Write driver	26
6.2	Write timing strategy	29
7	Control: Decoder, Clocking, Registers	29
7.1	Row decoder (4-to-16)	29

7.2	Non-overlapping clock generator	31
7.3	Input registers	32
8	Top-Level Integration and Validation Driver	32
8.1	The validation <code>.control</code> block	33
8.2	Top-level functional and timing result	33
8.3	Sustainable max CLK frequency sweep (<code>find_fmax.py</code>)	35
9	Critical-Path Timing Analysis	37
9.1	Read path (worst case)	37
9.2	Write path (worst case)	38
9.3	Access delay and operating point	39
10	Power Measurement	40
10.1	Methodology	40
10.2	Power decomposition (qualitative)	41
11	Figure of Merit	41
11.1	Numerical FOM	41
11.2	Sensitivity table	42
12	Validation Methodology and Summary	43
12.1	Test methodology	43
12.2	Test cases	43
12.3	Cross-reference of measured evidence	44
12.4	Spec compliance recap	44
13	Summary of Design Metrics	45
	Appendix A — Repository Map	46
	Appendix B — Electric Hierarchy Checklist	46
	Academic Integrity Statement	46

1 Executive Summary

This report presents the design, simulation, and characterization of a 16×4 SRAM memory (64 b total) in the PTM 22 nm High-Performance CMOS process at $V_{DD} = 1$ V. The design is optimized end-to-end for the project Figure of Merit

$$\text{FOM} = 60 \cdot \text{BitcellArea} \cdot \text{Power} \cdot \text{Delay}^2,$$

in which the quadratic dependence on delay drives every architectural choice. The final design uses a 16-row by 4-column array with *no* column multiplexer (all four bits of an addressed word are read or written in parallel), a six-transistor (6T) bitcell with asymmetric sizing CR=2, PR=1, a clocked StrongARM sense amplifier, a self-timed sense-amp enable generated by a replica bitline column, and a two-phase non-overlapping clock scheme that gates precharge, write driver, and wordline activity into non-conflicting time windows.

1.1 Required vs. Delivered (headline)

The project specification ([proj2.pdf](#)) defines a small set of hard requirements for a passing design: a 16×4 SRAM, single CLK, $V_{DD} \leq 1$ V, and a maximum operating frequency ≥ 500 MHz. Figure 1 places the spec value alongside the implemented value for the headline metrics in five panels (two frequency views on top so equivalent latency f_{eq} does not share an axis with sustainable f_{max}); the precision table immediately after lists the full set of measured numbers and points each row at its source artifact in the repository.




Figure 1: Required-vs-Delivered headline (horizontal bars): **top** sustainable f_{max} and equivalent f_{eq} on a dedicated 0–10 GHz axis; **bottom** delay, read-disturb, and DFF panels. For **frequency**, a longer bar to the right is better; for **delay**, **voltage disturb**, and **DFF delay**, the delivered (blue) bar should be shorter than the budget/threshold (gray). T_{min} and sweep context: Table 1 and §. All bar-chart numbers match Table 1.

Spec / Metric	Required	Delivered	Margin	Source
Capacity	16×4	16×4	—	top-level netlist
Supply V_{DD}	≤ 1.00 V	1.00 V	$1.00\times$	spice/top.spi
Sustainable max CLK frequency $f_{\max}$ (back-to-back W/W/R/R)	≥ 500 MHz	4.571 GHz	$9.14\times$	find_fmax.py sweep, §
Min sustainable CLK period $T_{\min}$	≤ 2 ns	218.75 ps	$9.14\times$	find_fmax.py sweep, §
Equivalent frequency $f_{eq} = 1/t_{\text{CLK} \rightarrow D_{out}}$ (single-edge latency)	—	9.037 GHz	—	Fig. 37
Top-level $t_{\text{CLK} \rightarrow D_{out}}$ (single read access)	< 2 ns (cycle)	110.65 ps	$18.07\times$	Fig. 37
Bitcell write $1 \rightarrow 0$ time	< 2 ns	9.87 ps	$203\times$	Test 1 (§)
Bitcell write $0 \rightarrow 1$ time	< 2 ns	21.8 ps	$91.7\times$	Test 2 (§)
Read disturb $\overline{Q}_{\max}$ (store 0)	≤ 0.4 V	95.4 mV	$4.19\times$	Test 3 (§)
Read disturb $\overline{Q}_{\max}$ (store 1)	≤ 0.4 V	95.4 mV	$4.19\times$	Test 4 (§)
Retention drop over 5 ns	≤ 50 mV	9.5 μ V	$> 5000\times$	Test 5 (§)
DFF CLK-to-Q average	< 100 ps	67.0 ps	$1.49\times$	Fig. 13
Functional readback @ 5 ns (0x5)	PASS	0x5	PASS	Fig. 37
Functional readback @ 7 ns (0xA)	PASS	0xA	PASS	Fig. 37
Bitcell summed width	—	8 $W_{\min}$	—	§
Average power (top-level run)	—	21.367 μ W	—	Fig. 37

Table 1: Headline measurement table. “Required” is the literal spec constraint or the conservative threshold used to declare each test PASS/FAIL; “Delivered” is the value taken directly from the in-repository simulation evidence cited in the right-hand column. No synthetic numbers appear in this table.

Final Figure of Merit $\text{FOM} = 60 \cdot \text{BitcellArea} \cdot \text{Power} \cdot \text{Delay}^2 =$
 $60 \cdot 8 \cdot (2.1367 \times 10^{-5} \text{ W}) \cdot (110.65 \text{ ps})^2 \approx \mathbf{1.256} \times \mathbf{10^{-22}}.$

1.2 Map of the report

The remainder of the report is organized so each claim sits next to the artifact that supports it. Figure 2 sketches the reading order from architecture down to validation. § establishes the architectural decisions and the FOM sensitivity that justifies them. § documents the primitive cells (inverter, NAND, transmission gate, DFF) that every higher-level block instantiates. § presents the 6T bitcell with the verbatim SPICE .SUBCKT, sizing rationale, and the five cell-level validation tests with their measured numbers boxed inline. § documents the column slice end-to-end (precharge, write driver, StrongARM sense, output latch). § documents the replica column for self-timed SAE. § documents the row decoder, non-overlap clock generator, and input registers. § presents the top-level integration, the validation .control block, and the terminal output it produced. § presents the critical-path timing analysis. § presents the power-measurement methodology and result. § presents the FOM calculation and the sensitivity table. § presents the validation-test methodology and the final summary of design metrics.





Figure 2: Report flow: bottom-up presentation from primitives and bitcell through column and control to top-level evidence, then metrics.

2 Architecture and FOM Strategy

2.1 Reading the FOM

The FOM penalizes three quantities. Bitcell area appears linearly, power appears linearly, and delay appears quadratically. A straightforward implication is that a $2\times$ improvement in delay is worth a $4\times$ degradation in either area or power. Equally important is what the FOM does *not* penalize: it does not penalize periphery transistor count or area. The bitcell-area term is defined in the project specification as the sum of transistor widths *in the repeated memory cell in the memory core*, not the area of the row decoder, sense amplifiers, write drivers, or clock generation.

This asymmetry has a single sharp consequence. Any optimization that moves complexity *out* of the bitcell and *into* the periphery is FOM-positive (provided it does not cost delay or power), and any optimization that reduces delay is doubly rewarded. Figure 3 compresses the weighting and the bitcell-vs-periphery asymmetry into one picture.



Decoder, sense amps, write drivers, replica column, and clock generation *do not* enter the bitcell-area term; only the repeated 6T core width does (proj2.pdf definition).

Figure 3: Project FOM weights: linear in bitcell-area sum and power, quadratic in delay; periphery area is off-book for the area term.

2.2 Top-level array organization: 16×4, no column mux

The data word width is four bits and the depth is sixteen words. There are two natural organizations:

1. **16 rows × 4 columns**, no column mux, all four bits read or written in parallel. The 4-bit address $A[3:0]$ is consumed entirely by the row decoder.
2. **4 rows × 16 columns** (or some intermediate folding), with column multiplexing reducing the array aspect ratio. Some number of address bits drive the row decoder, the rest drive a 4-to-1 (or larger) column mux.

I selected option (1). The reasoning, in FOM terms, is direct:

- **Delay penalty of column muxing.** A column multiplexer inserts a pass-gate stage in series with the bitline-to-sense-amp path on every read. At 22 nm HP a sized pass-gate adds $\approx 30\text{--}50$ ps to the path. With Delay^2 in the FOM, even a 40 ps adder on a ~ 110 ps access-time path inflates the squared-delay term by $\sim 85\%$.
- **Bitline capacitance penalty of column muxing.** A muxed column shares its sense amplifier, so each sense-amp input sees the bitlines of multiple columns through the mux pass-gates. This increases effective C_{BL} .
- **No bitcell-area savings.** The bitcell is unchanged regardless of array organization; the bitcell-area metric is the same in both cases.
- **Periphery cost is free under the FOM.** Option (1) requires a 4-to-16 decoder where option (2) might use a 2-to-4 decoder plus a 4-to-1 mux. The total periphery transistor count is comparable, and in any case is not weighted by the FOM.

The 16-row aspect ratio also keeps the bitline length short (16 cells per BL), which directly reduces C_{BL} and therefore the BL discharge component of read delay. Figure 4 contrasts the implemented floorplan with a mux-heavy alternative.




Figure 4: Array organization: (a) this design reads all four bits in parallel with no column select in the BL path; (b) a folded/muxed array adds series devices and bitline loading (Sec. 2.2).

2.3 Subsystem-level decisions

The remaining major architectural choices, with one-line FOM justifications, are:

- **6T bitcell, not 8T or 10T.** An 8T cell adds two transistors per cell and increases summed bitcell width by $\sim 60\%$ for a delay benefit that is small at this array size. 6T is FOM-optimal here. (§3.1)
- **Asymmetric sizing** $CR = 2$, $PR = 1$. $CR = 2$ is the minimum cell ratio that gives positive read static-noise margin across worst-case PTM HP corners at $V_{DD} = 1\text{ V}$. $PR = 1$ is the maximum pull-up ratio that still allows the pass-gate to flip the cell during a write. (§3.2)
- **Clocked StrongARM sense amplifier.** A clocked regenerative latch resolves a 100 mV BL differential in $\sim 40\text{--}50\text{ ps}$ and consumes *zero* static current. A static current-mirror sense amp resolves the same differential in $\sim 150\text{ ps}$ and burns continuous DC current, which would catastrophically inflate the power term measured during the write-stress test. (§4.3)
- **Replica bitline for SAE generation.** A self-timed SAE derived from a dummy column tracks the real BL discharge across PVT corners. A hardcoded inverter-chain delay would either trigger SAE too early in slow corners (sense amplifies noise) or too late in fast corners (wasted delay). (§4.4)
- **Two-phase non-overlapping clock.** A clean separation of the precharge phase Φ_1 from the active phase Φ_2 eliminates contention between write drivers and precharge transistors, which would otherwise cause large short-circuit current spikes that dominate the power measurement. (§6.2)
- **Edge-triggered input registers on A , D_{in} , WE .** Mid-cycle changes on the inputs are filtered out, so a single read or write operation is associated with a single CLK edge as the spec requires. (§6.3)

2.4 Block diagram of the full memory

As shown in Figure 5 (full top-level view) and the fan-out closeup in Figure 6, the implemented hierarchy keeps all FOM-neutral complexity in periphery blocks (decoder, clocking, replica timing, sensing, and latching) while preserving the compact 6T core bitcell in the repeated array. Nine edge-triggered input DFFs (four address, four D_{in} , one WE) sit at the boundary; the registered $WE \cdot \Phi_2$ produces WD_EN through a $NAND2 + \text{inverter}$; the four `column_slice` blocks each

carry one bit of the word; and the `replica_column` sits on the array periphery sharing all sixteen wordlines.




Figure 5: Top-level 16×4 SRAM Electric schematic (`top{sch}`, exported from `spice/proj2.jelib`). Four `column_slice` blocks produce `Dout[3:0]`; nine input DFFs register `A[3:0]`, `Din[3:0]`, and `WE` on the rising edge of `CLK`; the `predecoder_4to16` drives sixteen wordlines shared between the array and the `replica_column`; the `clock-gen-2phase` block produces `PCH_b`, `WL_EN`, and Φ_2 from the single `CLK` input.




Figure 6: Top-level closeup: the four `column_slice` blocks share all sixteen wordlines $WL_0 \dots WL_{15}$ from the row decoder and produce $D_{out}[3:0]$ in parallel. There is no column multiplexer between the bitlines and the sense amplifier inside each slice, so all four bits of an addressed word are read or written simultaneously.

3 Primitives: building blocks the rest of the design instantiates

Before the bitcell, the column, the decoder, or the top-level netlist can do anything, every one of those blocks is a composition of five primitive cells. Each primitive is sized once for its target fanout and then reused throughout the hierarchy. The four that matter most for this design are the inverter (used everywhere as a buffer or skewed trip point), the NAND2/NAND3 (used in the predecoder, row decoder, clock generator, and write-enable AND), the transmission gate (used in the DFF and the output latch), and the D flip-flop (instantiated nine times at the input boundary as the master-slave register that turns mid-cycle input changes into a non-issue). Functional waveforms

for each are shown in the figures of this section; the only primitive whose *measured timing* feeds directly into a top-level number is the DFF, which is therefore characterized in detail.

3.1 Inverter (and its skewed variants)

The inverter primitive (Fig. 7) is parameterized by a single drive-strength multiplier. The unit cell uses $W_n = W_p = 1 W_{\min}$; $\times 2$, $\times 4$, $\times 5$, and $\times 8$ variants scale both transistors proportionally. Skewed variants (asymmetric W_n/W_p) appear in two places: the precharge buffer ($W_p = 16$, $W_n = 8$ for fast falling edge of $\overline{\text{PCH}}$), and the dummy-bitline trip inverter inside the replica column (skewed high so its trip point sits near $V_{DD} - 100 \text{ mV}$).




Figure 7: Inverter primitive (left) and SPICE transient verification of the $\times 4$ and $\times 8$ buffers (*media/waveforms/inv_x4.png*, *inv_x8.png*) used as drivers in the row-decoder output and PCH_b buffer respectively.

3.2 NAND2 and NAND3

The decoder fabric uses two NAND types: NAND2 in the 2-to-4 predecoders, and NAND3 in the per-row select-and-enable gate (the third input is WL_EN , folded into the final stage to remove one inverter level from the WL critical path). Both are sized with stack equalization in mind: the NAND2 uses $W_n=W_p=4$ (driver-strength matched to a single PMOS pulling against two NMOS in series), and the NAND3 uses $W_n=4$, $W_p=4$ accepting the slower fall through three stacked NMOS in exchange for a sharper rise that beats the WL load.




Figure 8: NAND2 schematic (left) and SPICE transient (*media/waveforms/nand2.png*, right). Used as the 2-to-4 predecoder gate.




Figure 9: NAND3 SPICE transient (*media/waveforms/nand3.png*). Used per row in the 4-to-16 decoder; third input is WL_EN, so wordline gating costs zero extra logic levels.

3.3 Transmission gate

The transmission gate is the bidirectional pass element used inside the master-slave DFF and the output latch. Its merit relative to a single NMOS pass-gate is that it passes both rails cleanly (no V_{th} drop), which is essential when the held node feeds an inverter trip point. Two sized variants are used: a $W_n=W_p=2$ unit cell for the DFF (`tgate2`), and a $W_n=W_p=6$ variant for the output latch (`tgate`) where the latched node drives the D_{out} pad buffer.





Figure 10: Transmission-gate schematic (left) and pass-through transient (*media/waveforms/tgate.png*, right). Confirms full rail-to-rail pass behavior used by the DFF master/slave stages and the output latch.

3.4 D flip-flop (transmission-gate master-slave)

The DFF (Fig. 11) is the only primitive whose measured timing appears directly in the headline table. It is built from four `tgate2` pass elements and four `invX2_2` ($W_n=2, W_p=4$) inverters in a standard transmission-gate master-slave topology, with an output booster inverter `inv_X5` ($W_n=12, W_p=24$) sized to drive the registered output's downstream fanout (the row decoder for address bits, the column-slice D_{in} pin for data bits). The DFF is instantiated nine times at the input boundary: four $A[3:0]$, four $D_{in}[3:0]$, and one WE .


```

=====
  COMPLETE DFF VALIDATION
=====
CLK-to-Q rise delay:
t_cq_rise = 4.202258e-11
CLK-to-Q fall delay:
t_cq_fall = 9.193923e-11

Average CLK-to-Q:
t_cq_avg = 6.698090e-11


Multi-cycle test (2nd rise):
t_cq_rise2 = 4.202258e-11

ELITE: Ultra-fast (<100ps average)

Rise/Fall Balance:
ratio = 4.570691e-01
EXCELLENT: Well-balanced (<1.5x ratio)
=====

```

Figure 13: NGSPICE `.measure` output for the DFF validation deck (`media/terminal/dfvalidation.png`). The numbers below are read directly from this terminal capture.



The average $t_{cq} = 67.0 \text{ ps}$ fits inside the cycle budget across the tested range, including the delay-derived equivalent operating point $f_{eq} \approx 9.037 \text{ GHz}$ ($t_{CLK \rightarrow D_{out}} \approx 110.65 \text{ ps}$); at the 500 MHz spec floor the DFF margin is $2 \text{ ns}/67 \text{ ps} \approx 30\times$.

4 Bitcell Design

4.1 Topology: 6T

The bitcell is the standard six-transistor SRAM cell: two cross-coupled CMOS inverters forming the bistable storage element, plus two NMOS access (pass-gate) transistors connecting the storage nodes to the differential bitlines BL and BLB when the wordline WL is asserted. The schematic, with device labels used throughout the report, is shown in Figure 14.




Figure 14: 6T SRAM bitcell. Two cross-coupled CMOS inverters form the bistable storage element (Q , $\bar{Q}$); two NMOS pass-gates connect Q and $\bar{Q}$ to BL and BLB when WL is asserted. Sizing: pull-down NMOS $W_{PD}=2W_{\min}$, pull-up PMOS $W_{PU}=1W_{\min}$, pass-gate NMOS $W_{PG}=1W_{\min}$, giving $CR=2$ and $PR=1$. Summed transistor width per cell is $8W_{\min}$, the value reported as the bitcell-area FOM term.

The verbatim .SUBCKT from `spice/top.spi` is reproduced below to make every cited width auditable; the entire memory instantiates this single subcircuit $16 \times 4 = 64$ times in the array plus 16 instances of `dummy_bitcell` in the replica column.

```
.SUBCKT proj2__sram_6t_cell bl blb wl
** GLOBAL gnd
** GLOBAL vdd
Mnmos@3 qb q gnd gnd N L=2 W=2 * pull-down (CR=2)
Mnmos@4 q qb gnd gnd N L=2 W=2 * pull-down (CR=2)
Mnmos@5 bl wl q gnd N L=2 W=1 * pass-gate
Mnmos@6 blb wl qb gnd N L=2 W=1 * pass-gate
Mpmos@0 qb q vdd vdd P L=2 W=1 * pull-up (PR=1)
Mpmos@1 q qb vdd vdd P L=2 W=1 * pull-up (PR=1)
.ENDS proj2__sram_6t_cell
```

4.2 Sizing and rationale

The summed transistor width is the term the FOM cares about, so each transistor must justify its width. Widths are expressed in units of $W_{\min}$ (the technology minimum drawn width as defined in the Electric tech file).

Device	Width ($W_{\min}$)	Function
$M_{PD,L}, M_{PD,R}$ (NMOS)	2	Pull-down in cross-coupled inverter
$M_{PU,L}, M_{PU,R}$ (PMOS)	1	Pull-up in cross-coupled inverter
$M_{PG,L}, M_{PG,R}$ (NMOS)	1	Access (pass-gate) to BL/BLB
Total summed W	8	Bitcell area metric

Table 2: Bitcell sizing. Cell ratio $CR = W_{PD}/W_{PG} = 2$. Pull-up ratio $PR = W_{PU}/W_{PG} = 1$.

The two key ratios are the cell ratio (CR) and pull-up ratio (PR). They govern, respectively, read stability and writeability.

4.2.1 Cell ratio: read static-noise margin

When WL rises during a read, the storage node holding “0” is pulled *up* through the pass-gate by the precharged bitline. This is the read-disturb mechanism: if the pull-down cannot sink the pass-gate current fast enough, the storage node climbs above the trip point of the opposing inverter and the cell is destructively flipped. The voltage divider in the limit of static analysis is

$$V_{Q,\text{disturb}} \approx V_{DD} \cdot \frac{R_{PD}}{R_{PD} + R_{PG}} \approx V_{DD} \cdot \frac{1}{1 + CR \cdot (\mu_{PG}/\mu_{PD})},$$

where the mobility-correction factor is ≈ 1 for two NMOS devices in the same technology. With $CR = 1$, this gives $V_{Q,\text{disturb}} \approx 0.50$ V, which sits exactly at the inverter trip point of a balanced CMOS inverter at $V_{DD} = 1$ V. The read SNM in this case is essentially zero, and across the HP fast-NMOS slow-PMOS corner the cell will read-disturb.

With $CR = 2$, $V_{Q,\text{disturb}} \approx 0.33$ V, which is comfortably below any plausible inverter trip point. This prediction is verified directly in Test 3 and Test 4 of the cell-level validation (§), where the maximum read-disturbed storage voltage measured under a sustained WL pulse is 95.4 mV — a 4.2 $\times$ margin below the 0.4 V threshold.

4.2.2 Pull-up ratio: write margin

During a write, one bitline (say BLB) is pulled to ground by the write driver. The storage node holding “1” must be pulled *down* through its pass-gate, against the resistance of the PMOS pull-up that is currently sourcing current to keep that node high. If the pull-up is too strong relative to the pass-gate, the cell will not flip.

The dominant constraint at this node is that the NMOS pass-gate (W_{PG}) must out-pull the PMOS pull-up (W_{PU}). NMOS in 22 nm HP has roughly 2 $\times$ the on-current of PMOS at equal W , so even $PR = W_{PU}/W_{PG} = 1$ already gives the NMOS a 2 $\times$ strength advantage; $PR > 1.5$ pushes us into territory where write failures appear at slow-NMOS / fast-PMOS corners. The chosen sizing’s write behavior is verified directly in Test 1 and Test 2 of the cell-level validation (§), with measured write times 9.87 ps ($1 \rightarrow 0$) and 21.8 ps ($0 \rightarrow 1$) — both more than two orders of magnitude faster than the 2 ns cycle budget at the spec floor.

4.2.3 Why not minimum everything?

A bitcell sized 1-1-1 (all transistors at $W_{\min}$) has summed $W = 6$ instead of 8, a 25% reduction in the bitcell area term. The temptation is real but is the wrong choice: at $CR = 1$ the read

SNM is essentially zero at $V_{DD} = 1$ V across HP corners, and at least one of the validation tests (read-disturb stress) fails. The 6T sizing 2-1-1 is the FOM-minimal sizing that still yields a *working* memory. We do not give back functionality for a 25 % area shave that the FOM would partially recover anyway, given that the broken design has effectively undefined power and delay.

Figure 15 gives the physical row-level view that this organization implies: one asserted wordline activates four columns in parallel with no column-select stage in between.




Figure 15: One row of four 6T bitcells stamped at the array pitch, sharing WL and presenting per-cell loads onto four independent BL/BLB column pairs. Because the architecture has no column multiplexer, all four bits of an addressed word are accessed in parallel through this row, and the wordline load is exactly four pass-gate gates plus the dummy pass-gate from the replica column.

4.3 Bitline capacitance

The bitline capacitance is the dominant load on the read path and is needed for both delay and power estimates.

Pre-layout estimate. With 16 cells per BL plus the replica column’s bitcell, two PG drains per cell, and short M2 wire at the array pitch:

- $C_{PG,drain} \approx 0.25$ fF per cell at $W_{PG} = 1$.
- $C_{wire} \approx 0.15$ fF per cell pitch on M2.

Per-bitline pre-layout capacitance is $C_{BL} \approx 16 \times 0.40 + 1$ (SA input) ≈ 7.4 fF.


Post-extraction estimate. Diffusion fringe, M1/M2 coupling, and sense-amp input cap roughly double this. The post-extraction target used for the integrated delay budget is $C_{BL} \approx 13$ fF. This is the value used for the timing analysis in §. Floorplan rules (M1 wordline / M2 bitline, perpendicular routing, BL-VDD-BLB shielding where the M2 track count permits) keep the post-extraction value within $\sim 30\%$ of the pre-layout estimate; without these rules it can climb past 18 fF.

4.4 Cell-level validation simulations: Tests 1–5

Five unit tests exercise the bitcell with ideal PWL drive on WL, BL, and BLB. Each test prints its timing/voltage measurement to a terminal capture (`.measure ... print` sequence under `.control`); each measurement is shown next to its waveform below. All five PASS criteria are met by a margin of at least $4\times$ on every metric.

4.4.1 Test 1 — Write 1→0

Initialize $Q = 1$, drive $BL = 0$, $BLB = V_{DD}$, pulse WL high. Measure the elapsed time from WL rise to Q falling through the inverter trip point.




```
=====
TEST 1: WRITE 1->0
=====
Write time (WL rise -> Q fall):
t_write_1to0 = 9.869044e-12
PASS: Write-ability excellent (<50ps)
=====
```

Figure 16: Test 1. Left: SPICE transient (*media/waveforms/test1_write_1to0.png*). Right: NGSPICE `.measure` terminal output (*media/terminal/test1-sramwrite1-0.png*).

Test 1: $t_{\text{write},1\rightarrow0} : 9.869 \text{ ps}$	PASS ($< 50 \text{ ps}$).
---	-----------------------------

4.4.2 Test 2 — Write 0→1

Mirror of Test 1 with the initial cell state reversed and the bitline drive flipped, exercising the writeability of the opposite storage node.



```

=====
TEST 2: WRITE 0->1
=====
Write time (WL rise -> Q rise):
t_write_0to1 = 2.176203e-11
PASS: Bidirectional write excellent (<50ps)
=====


```

Figure 17: Test 2. Left: SPICE transient (*media/waveforms/test2_write_0to1.png*). Right: NGSPICE *.measure* terminal output (*media/terminal/test2-sram-write0->1.png*).

Test 2: $t_{\text{write},0 \rightarrow 1}$: 21.762 ps	PASS (< 50 ps).
---	-----------------

4.4.3 Test 3 — Read stability storing 0

With $Q = 0$, $\overline{Q} = 1$, precharge BL = BLB = V_{DD} and pulse WL. The read-disturbed storage node is Q , which must not climb past the opposing inverter's trip point.



```

=====
TEST 3: READ STABILITY (Storing 0)
=====
Max Q voltage during read:
v_q_max = 9.535303e-02
PASS: Cell is read-stable (Q < 0.3V)
=====


```

Figure 18: Test 3. Left: SPICE transient (*media/waveforms/test3_read_stable_0.png*). Right: NGSPICE *.measure* terminal output (*media/terminal/test3-read-stability.png*).

Test 3: max V_Q during read : 95.4 mV	PASS (< 300 mV; $\sim 4.2\times$ below 0.4 V destructive threshold).
--	--

4.4.4 Test 4 — Read stability storing 1

Mirror of Test 3 with the cell state reversed; the read-disturbed node is now $\overline{Q}$.



```

=====
TEST 4: READ STABILITY (Storing 1)
=====
Max QB voltage during read:
v_qb_max = 9.535303e-02
PASS: Cell is read-stable (QB < 0.3V)
=====


```

Figure 19: Test 4. Left: SPICE transient (*media/waveforms/test4_read_stable_1.png*). Right: NGSPICE *.measure* terminal output (*media/terminal/test4-read-stability.png*).

Test 4: max V_Q during read: 95.4 mV	PASS (< 300 mV).
---	------------------

4.4.5 Test 5 — Data retention (hold)

With $WL = 0$ for 5 ns, monitor the storage node for drift. This is the cell's intrinsic bistability check; any non-trivial leakage imbalance shows up here.



```

=====
TEST 5: DATA RETENTION
=====
Min Q voltage (should stay ~1.0V):
v_q_min = 9.999905e-01
Final Q voltage at t=5ns:
v_q_end = 9.999905e-01
PASS: Data retained perfectly
=====

```

Figure 20: Test 5. Left: SPICE transient (*media/waveforms/test5_retention.png*). Right: NGSPICE *.measure* terminal output (*media/terminal/test5-data-retention sram.png*). Final V_Q at $t = 5$ ns equals the minimum V_Q to within numerical precision; drift is below the simulator's reporting resolution.

Test 5: min V_Q over 5 ns: 0.9999905 V	PASS (drift < 10 μ V over 5 ns).
---	--------------------------------------

5 Read Path (Column slice, read direction)

The read operation, in time order, is: precharge BL/BLB to V_{DD} during Φ_1 ; raise WL on Φ_2 ; the addressed bitcell discharges one of (BL, BLB) through its pull-down and pass-gate; the replica column discharges its dummy BL in parallel and triggers SAE when $\Delta V_{BL} \approx 100$ mV; the StrongARM sense amplifier regenerates a full-rail logic level from this differential; the output latch captures it before the next precharge phase.




Figure 21: Single-column SPICE testbench used throughout §–§, integrating every periphery block this report describes: precharge and equalize at the top, the tristate write driver, sixteen 6T bitcells stacked along BL/BLB, the StrongARM sense amplifier, and the output latch. All read-path and write-path delay numbers in § are measured on this column with the post-extraction bitline load $C_{BL} \approx 13$ fF.

5.1 Precharge and equalize

Each column has two PMOS precharge devices (one to BL, one to BLB) and one PMOS equalize device shorting BL to BLB. All three are gated by PCH_b (active low). Sizing:

- Precharge PMOS, $W = 4W_{\min}$ each. Sized to recover BL/BLB from full-swing GND (post-write) to $V_{DD} - 50\text{ mV}$ in $\leq 100\text{ ps}$; this fits inside Φ_1 .
- Equalize PMOS, $W = 2W_{\min}$. Equalization time $\ll$ precharge time because it operates on the small mismatch rather than the full rail.

The PCH_b control signal is driven by an inverter sized $W_p = 16$, $W_n = 8$ to give a sub-30 ps rise time. A slow rise on PCH_b creates a window in which the precharge devices and the write driver are both partially on, a short-circuit current path that inflates the power measurement. Fast PCH_b edges close this window.





Figure 22: Per-column precharge and equalize block. Two PMOS precharge devices ($W=4W_{\min}$) tie BL and BLB to V_{DD} ; one PMOS equalize device ($W=2W_{\min}$) shorts BL to BLB. All three gates are driven by PCH_b through a strong inverter ($W_p=16$, $W_n=8$) so the precharge edges close in $<30\text{ ps}$, eliminating the short-circuit window between precharge and write driver.



```

=====
PRECHARGE & EQUALIZATION RESULTS
=====
Precharge time (BL 0→0.95V):
t_pch_bl = 1.798144e-10
Precharge time (BLB settling):
t_pch_blb = 1.570471e-10
Equalization time:
t_equalize = 1.587612e-10

PASS: Precharge fast (<200ps)
=====

```

Figure 23: Precharge / equalize evidence. Left: SPICE transient (*media/waveforms/precharge_test.png*) showing both BL and BLB recovering to V_{DD} after a write-induced full-rail excursion. Right: NGSPICE *.measure* terminal output (*media/terminal/prechargeeq_results.png*) reporting the measured times for both bit-lines and the equalize.

$t_{\text{pch,BL}}$ (BL 0→0.95 V) : 179.8 ps	$t_{\text{pch,BLB}}$: 157.0 ps	t_{equalize} : 158.8 ps
PASS (< 200 ps).		

5.2 Bitcell BL discharge

The selected cell, with $Q = 0$, $\overline{Q} = 1$, WL asserted, will pull BL down through the series combination of the left pass-gate and the left pull-down. (The other side, BLB, is held near V_{DD} by the right pass-gate's source, since $\overline{Q} = 1$.) The discharge time to a 100 mV differential is

$$t_{BL,100\text{mV}} \approx \frac{C_{BL} \cdot 100 \text{ mV}}{I_{\text{cell}}},$$

where I_{cell} is the cell read current set by the series combination of M_{PG} and M_{PD} . With the chosen sizes at the typical corner, $I_{\text{cell}} \approx 35 \mu\text{A}$, giving $t_{BL,100\text{mV}} \approx (13 \text{ fF} \cdot 0.1 \text{ V}) / 35 \mu\text{A} \approx 37 \text{ ps}$ post-extraction.

5.3 StrongARM sense amplifier

The sense amplifier is a clocked StrongARM regenerative latch (Figure 24). When $\text{SAE} = 0$, the tail NMOS is off and internal nodes precharge to V_{DD} via the cross-coupled latch's PMOS reset transistors. When SAE rises, the tail turns on and the input differential pair (whose gates are tied to BL and BLB) discharges the internal nodes at slightly different rates determined by the input differential. Once one internal node falls past $V_{DD} - V_{tp}$, the cross-coupled latch regenerates to full rails.

Sizing (per sense amp, four total):

- Input pair NMOS: $W = 4$ each.
- Cross-coupled latch NMOS: $W = 4$ each.
- Cross-coupled latch PMOS: $W = 4$ each.
- Tail NMOS: $W = 8$.
- Reset PMOS pair (precharge of internal nodes): $W = 2$ each.




Figure 24: Clocked StrongARM regenerative sense amplifier. Differential NMOS input pair ($W=4W_{\min}$, gates on BL/BLB) discharges the internal nodes of a cross-coupled CMOS latch ($W_n=W_p=4W_{\min}$) at slightly different rates set by the input differential; once one internal node falls below $V_{DD}-V_{tp}$, the latch regenerates to full rails. Tail NMOS ($W=8W_{\min}$) gated by SAE provides clock enable; PMOS reset devices ($W=2W_{\min}$) precharge the internal nodes while SAE is low. Resolves a 100 mV input differential in ~ 45 ps with zero static current.

The verbatim .SUBCKT proj2__sense_amp from spice/top.spi is reproduced below. The widths shown match the sizing list above; the topology is a textbook StrongARM with PMOS reset (Mpmos@0,3) and a single tail NMOS (Mnmos@4) gated by SAE.

```
.SUBCKT proj2__sense_amp bl blb out outb SAE
** GLOBAL gnd
** GLOBAL vdd
Mnmos@0 out outb node2 gnd N L=2 W=6 * cross-coupled latch NMOS
Mnmos@1 outb out node3 gnd N L=2 W=6 * cross-coupled latch NMOS
Mnmos@2 node2 bl node1 gnd N L=2 W=4 * input pair NMOS (BL gate)
Mnmos@3 node3 blb node1 gnd N L=2 W=4 * input pair NMOS (BLB gate)
Mnmos@4 node1 SAE gnd gnd N L=2 W=8 * tail NMOS (SAE enable)
```



```

Mpmos@0 vdd SAE out vdd P L=2 W=2 * reset PMOS (out)
Mpmos@1 vdd outb out vdd P L=2 W=6 * cross-coupled latch PMOS
Mpmos@2 vdd out outb vdd P L=2 W=6 * cross-coupled latch PMOS
Mpmos@3 vdd SAE outb vdd P L=2 W=2 * reset PMOS (outb)
.ENDS proj2__sense_amp

```

The decision-time-vs-input-differential characteristic is the critical sizing curve and is shown in Figure 25. At $\Delta V_{in} = 100$ mV the SA resolves in roughly 45 ps. The replica column (§) is tuned to fire SAE when the real BLs are at roughly this differential.




Figure 25: StrongARM sense-amplifier transient waveform evidence at $V_{DD} = 1$ V, showing rapid regeneration once SAE is asserted and a differential is present at BL/BLB.

Offset note. StrongARM input-referred offset due to threshold mismatch on the input pair is approximately $\sigma_{V_{th}}/\sqrt{WL} \approx 15\text{--}25$ mV (1σ) at $W = 4W_{\min}$. The 100 mV trigger gives a $\sim 4\text{--}6\sigma$ margin per SA. With four SAs, the worst sets the delay; a Monte Carlo analysis is outside the scope of this assignment, but the design has sufficient margin for the deterministic corners simulated.

5.4 Replica column for self-timed SAE

5.4.1 Motivation

A hardcoded SAE-after-WL delay (e.g. a 4-inverter chain) is fragile. Across PVT corners, the inverter chain’s delay scales with $\mu \cdot V_{DD}/V_{th}^2$, while the BL discharge time scales with $C_{BL} \cdot \Delta V/I_{\text{cell}}$ where I_{cell} itself scales with $\mu \cdot (V_{GS} - V_{th})^2$. The two do not track identically. In a slow corner, BL discharge slows down more than the inverter chain does, so the SAE chain fires *before* the BLs have developed sufficient differential and the SA amplifies noise. In a fast corner, the SAE fires after the BLs have already developed much more than 100 mV, wasting delay and burning extra power on BL

swing.

5.4.2 Replica circuit

The replica column is one extra column (17 columns total in the physical layout) of dummy bitcells. Each dummy bitcell is identical to the regular bitcell except that the storage node Q is hardwired to GND, so $\bar{Q}$ is hardwired to V_{DD} . This guarantees that whenever any wordline rises, the dummy column's BL discharges exactly as if reading a "0". The wordlines are shared with the real array, so no special timing is required.

The dummy BL drives a CMOS inverter sized so that its trip point sits near $V_{DD} - 100\text{ mV}$, i.e. it fires when the dummy BL has discharged by roughly the same amount as the real worst-case BL needs to discharge for the SA to decide reliably. A short tunable delay chain (2–3 inverters) follows, allowing the SAE timing to be empirically dialed in during simulation. The output is SAE.

The hard-tied storage-node implementation used in the dummy column is shown in Figure 26; the full sensing-timing path that converts this tracked BL discharge into SAE is shown in Figure 27.




Figure 26: Dummy bitcell instance from the replica column. Layout and device sizing are identical to the production 6T bitcell, but the storage node Q is hardwired to GND (and $\bar{Q}$ to V_{DD}). On every wordline rise the dummy cell unconditionally discharges its BL through the matched pass-gate / pull-down stack, so a single dummy column tracks the real array's worst-case BL discharge across PVT corners.




Figure 27: Replica/dummy column for self-timed SAE generation. Sixteen dummy bitcells with Q hardwired to GND share all sixteen wordlines with the real array; the dummy BL is loaded identically to a real BL so its discharge tracks $C_{BL} \cdot \Delta V / I_{\text{cell}}$ across all PVT corners. The dummy BL feeds a CMOS inverter skewed so that its trip point sits near $V_{DD} - 100 \text{ mV}$, followed by a short (≤ 3 -stage) tunable inverter delay; the output is SAE, which fires at the moment the real bitlines have developed the $\sim 100 \text{ mV}$ differential the StrongARM SA needs to decide.

5.4.3 Robustness

Because the dummy BL and real BLs are both discharged by NMOS pull-down stacks of identical sizing, fed from the same wordlines, into matched M2 wire loads at the same array pitch, the dummy-BL discharge tracks the real-BL discharge across PVT. SAE fires at a roughly constant fraction of ΔV_{BL} regardless of corner. This is the canonical SRAM self-timed sensing scheme.

5.4.4 Bitcell-area accounting note

The spec defines the bitcell-area metric as “the sum of the total transistor width for the repeated memory cell in the memory core.” The dummy column contains repeated memory cells but is part of the sense-control periphery, not the storage core; its contents are hardwired and not addressable. The block diagram (§2.4) classifies the replica column under sense control. The bitcell-area metric is reported as the per-cell summed width times one (i.e. $8W_{\min}$), not multiplied by 64 or 68; the spec asks for the bitcell width, not the array width.

5.5 Output latch

After SAE fires and the StrongARM resolves, the SA outputs are captured by a transparent-low latch (transparent during Φ_2 , holding during Φ_1). This isolates D_{out} from the bitlines during the next precharge phase, when BL/BLB are pulled back to V_{DD} and the SA returns to its reset state. Without this latch, D_{out} would glitch back during precharge.

The latch is a standard pass-gate latch: a transmission gate gated by Φ_2 feeds a back-to-back inverter pair sized for a fanout-of-1 load on the output pad. The output buffer is a tapered inverter pair ($W = 4$ then $W = 8$) sized to drive a 5 fF nominal output load with a ~ 30 ps delay.




Figure 28: Per-column output latch and pad driver. The transmission gate is enabled by Φ_2 (transparent during the active phase, opaque during Φ_1 precharge), feeding a back-to-back inverter pair that holds the captured SA output. A tapered output buffer ($W=4W_{\min} \rightarrow W=8W_{\min}$) drives the 5fF nominal D_{out} pad load with a ~ 30 ps delay, and isolates D_{out} from BL/BLB during the next precharge so no glitch propagates.

6 Write Path (Column slice, write direction)

6.1 Write driver

Each column has one tristate write driver. When $WD_EN = 1$ (i.e. $WE = 1$ AND $\Phi_2 = 1$), the write driver drives BL and BLB to D_{in} and $\overline{D_{in}}$ respectively at full rails. When $WD_EN = 0$, both BL and BLB outputs are high-Z so they do not fight the precharge.

The write driver is implemented as two tristated inverters, one driving BL with input D_{in} and one driving BLB with input $\overline{D_{in}}$. Sizing per inverter: NMOS $W = 8$, PMOS $W = 16$ (matched current). The high-Z control is implemented by gating the pull-up branch with WD_EN_b and the pull-down branch with WD_EN via stacked enable transistors of the same width.




Figure 29: Tristate inverter primitive used twice per column inside the write driver. Stacked PMOS enable on top (gated by WD_EN_b) and stacked NMOS enable on the bottom (gated by WD_EN) put the output into a clean high-Z when the driver is disabled, so it never contends with the precharge devices. Sized at $W_n=8W_{\min}$, $W_p=16W_{\min}$ on the data branch.





Figure 30: Per-column tristate write driver. Two tristate inverters drive BL and BLB from D_{in} and $\overline{D_{in}}$ respectively at full rails when $WD_EN=1$ ($WE \cdot \Phi_2$); stacked enable devices hold both BL and BLB at high-Z otherwise so the precharge is never contended. Output inverters sized $W_n=8W_{min}$, $W_p=16W_{min}$ for matched current, comfortably stronger than the $W=4W_{min}$ precharge devices and Φ_2 -gated to ensure precharge has fully released before drive begins.



```

=====
WRITE DRIVER (13fF Bitline Load)
=====
Write 1 delay (din↑ → bl↑):
t_write_rise = 1.673257e-10
Write 0 delay (din↓ → bl↓):
t_write_fall = 1.429166e-10
=====

```

Figure 31: Write driver evidence into a 13fF bitline load. Left: SPICE transient (*media/waveforms/write_driver.png*). Right: NGSPICE *.measure* terminal output (*media/terminal/write-driver-load.png*).

$t_{\text{write,rise}} (D_{in} \uparrow \rightarrow BL \uparrow) : 167.3 \text{ ps}$	$t_{\text{write,fall}} (D_{in} \downarrow \rightarrow BL \downarrow) : 142.9 \text{ ps}$	PASS: WD
comfortably overpowers $W=4W_{min}$ precharge devices.		

6.2 Write timing strategy

The naive write sequence “raise WL, then drive BL/BLB” wastes delay: the cell flip cannot start until BL/BLB have reached final rails, and BL/BLB take time to slew because C_{BL} is large.

The chosen sequence overlaps these:

1. During Φ_1 (precharge phase), BL and BLB are precharged to V_{DD} .
2. Mid- Φ_1 , if $WE = 1$, the write driver is enabled (**not yet**—this would conflict with precharge). Instead, WD_EN is gated by Φ_2 so the write driver fires only after precharge has released.
3. At $\Phi_1 \rightarrow \Phi_2$ edge: precharge releases; write driver fires; WL rises.
4. BL/BLB slew toward their written values; the selected bitcell flips as soon as the differential is sufficient.




Figure 32: Write timing (qualitative): precharge lives in Φ_1 ; write driver and wordline activity are gated into Φ_2 , separated by non-overlap from precharge so BL/BLB are not fought.

The write driver, sized at $W_n = 8$, $W_p = 16$, is significantly stronger than the precharge devices ($W = 4$). The non-overlap between Φ_1 (precharge release) and Φ_2 (WD enable) ensures these never both drive at the same time.

Critical write timing requirement. The intended cell must see its WL rise *after* BL/BLB have crossed their respective inverter trip points, so that the cell flips *toward* the intended value rather than oscillating. This is automatically satisfied because (a) WL rise is gated by WL_EN which is itself gated by Φ_2 , (b) WD enable is gated by Φ_2 , and (c) the WL driver path includes the row decoder (4-stage critical path ~ 110 ps) while the WD path is one stage; so BL/BLB lead WL by the decoder delay. Verified in simulation (§).

Half-select / unselected-row safety. Bitcells in unselected rows have $WL = 0$ and so are isolated from BL/BLB regardless of what the write drivers are doing. There is no half-select disturbance in this architecture because the four columns of the selected word are all written simultaneously (no column mux), so either an entire word is being written or none.

7 Control: Decoder, Clocking, Registers

7.1 Row decoder (4-to-16)

The row decoder is implemented as a 2-stage NAND-NAND structure:

- Stage 1: two 2-to-4 predecoders. The first decodes $\{A_3, A_2\}$, the second decodes $\{A_1, A_0\}$. Each predecoder output is one of $\{\overline{A_i A_j}, \overline{A_i} \overline{A_j}, \overline{A_i} A_j, \overline{A_i} \overline{A_j}\}$, produced by a NAND2 followed by an inverter (so the predecoder outputs are active-high decoded lines).

- Stage 2: per row, a NAND2 of one line from each predecoder, followed by an inverter that drives the wordline through an WL_EN gating AND.

The WL_EN gate is folded into the final stage as a NAND3 (predecode A , predecode B , WL_EN) followed by an inverter, eliminating one logic level.

Sizing: input-stage NAND2 with $W_n = 2$, $W_p = 2$. Predecoder output inverter $W_n = 2$, $W_p = 4$ (driving 4 row NAND3 inputs each). Per-row NAND3 with $W_n = 3$, $W_p = 2$. Final WL-driver inverter $W_n = 4$, $W_p = 8$, driving the wordline load (4 pass-gate gates plus replica column dummy pass-gate plus M1 wire) which is approximately 3 fF post-extraction. WL rise time ≈ 20 ps.




Figure 33: 2-to-4 predecoder used twice in stage 1 of the row decoder (one decoding $\{A_3, A_2\}$, one decoding $\{A_1, A_0\}$). Each of the four outputs is a NAND2 of one address bit and one address-bit complement, followed by an output inverter so the predecoder lines are active-high. Sizing: NAND2 $W_n=W_p=2W_{\min}$; output inverter $W_n=2W_{\min}$, $W_p=4W_{\min}$ to drive the four row NAND3 gate inputs each predecoder line fans out to.




Figure 34: 4-to-16 row decoder, two-stage NAND-NAND topology. Two 2-to-4 predecoders ($\{A_3, A_2\}$ and $\{A_1, A_0\}$, NAND2+INV per output) feed sixteen per-row NAND3 gates, each fanning in one line from each predecoder plus WL_EN, so the wordline-enable gate is folded into the final logic stage and one inverter level is removed from the critical path. WL driver inverter sized $W_n=4W_{\min}$, $W_p=8W_{\min}$ to drive the $\sim 3\text{fF}$ wordline load (4 pass-gate gates + replica dummy gate + M1 wire) with a $\sim 20\text{ps}$ rise time.

7.2 Non-overlapping clock generator

The single CLK input is converted to two non-overlapping phases Φ_1 and Φ_2 using the standard cross-coupled NAND structure with a delay chain in each feedback path. The non-overlap gap is sized to $\approx 30\text{--}40\text{ps}$ at typical corner. Note that *this gap grows in slow corners and shrinks in fast corners*; the design constraint is to ensure the gap remains positive at the fastest corner the design must operate at.

The clock generator produces five derived signals from Φ_1 and Φ_2 :

- PCH_b: $\overline{\Phi_1}$ buffered through a strong inverter ($W_p = 16$, $W_n = 8$) for fast precharge edges.
- WL_EN: Φ_2 buffered.

- WD_EN: $\Phi_2 \cdot WE$ (gated AND with the registered write-enable).
- Address-register clock: CLK directly (registers capture on CLK rising edge).
- SAE: *not* generated here. SAE is self-timed by the replica column.




Figure 35: Two-phase non-overlapping clock generator. Cross-coupled NAND structure with a delay chain in each feedback path produces Φ_1 (precharge phase) and Φ_2 (active phase) with a $\sim 30\text{--}40\text{ ps}$ non-overlap gap at the typical corner. Output buffers tap the derived control signals: PCH_b ($W_p=16 W_{\min}$, $W_n=8 W_{\min}$ for sub-30 ps precharge edges), WL_EN (buffered Φ_2), and WD_EN (gated AND of Φ_2 with the registered WE).

7.3 Input registers

Address $A[3:0]$, data input $D_{in}[3:0]$, and write enable WE are all captured by positive-edge-triggered D flip-flops (the `proj2__dff` primitive characterized in § with measured $t_{cq,avg} = 67.0\text{ ps}$). Nine identical instances are instantiated at the input boundary: four for the address, four for D_{in} , and one for WE . The address DFFs feed the predecoder, the data DFFs feed the column-slice D_{in} pin of each `column_slice`, and the registered WE AND-ed with Φ_2 becomes WD_EN (`Xnand2@0 + Xinv_X4@0` in the top-level netlist). Mid-cycle changes on any input pin are filtered out, so a single read or write is associated with exactly one CLK rising edge as the spec requires.

8 Top-Level Integration and Validation Driver

This section ties the previous seven together: every primitive characterized in §, every cell-level result from §, every column-level result from §–§, and every control block from § is instantiated exactly once in `spice/top.spi` (the schematic exports of which are Figs. 5 and 6). The top-level netlist instantiates: nine `proj2__dff` input registers, one `proj2__clock-gen-2phase`, one `proj2__predecoder_4to16`, four `proj2__column_slice` (each containing the precharge-eq, sixteen

6T cells, the write driver, the StrongARM SA, and the output latch), one `proj2__replica_column`, plus the `nand2+inv_X4` that produce `WD_EN`.

8.1 The validation `.control` block

The terminal output reproduced as Fig. 37 is generated by the `.control` block embedded in `spice/top.spi`. The block runs a 12 ns transient at $V_{DD} = 1$ V with a 500 MHz CLK, drives PWL inputs that write `0x5` to address 0 and `0xA` to address 1 in the first 4 ns, deasserts `WE` to enter the read window, samples D_{out} at $t = 5$ ns and $t = 7$ ns, and computes hex-equivalent readbacks. The CLK-to- D_{out} delay measurement uses the `trig/targ` construct on the third CLK rising edge of the read window; the average V_{DD} supply current is taken over the full 12 ns window. Figure 36 maps the same sequence on a 12 ns axis. The verbatim excerpt is reproduced below.




Figure 36: Top-level validation transient (not to scale on CLK edges): writes in the first ~ 4 ns, then readback samples at 5 ns and 7 ns on the extracted deck (§).

```
.control
tran 10p 12n
* Readback sampling points
meas tran D0_at_5ns find v(Dout_0) at=5n
meas tran D1_at_5ns find v(Dout_1) at=5n
meas tran D2_at_5ns find v(Dout_2) at=5n
meas tran D3_at_5ns find v(Dout_3) at=5n
meas tran D0_at_7ns find v(Dout_0) at=7n
meas tran D1_at_7ns find v(Dout_1) at=7n
meas tran D2_at_7ns find v(Dout_2) at=7n
meas tran D3_at_7ns find v(Dout_3) at=7n
* Timing metric: CLK-to-Dout (read of address 0)
meas tran t_clk_to_dout trig v(CLK) val=0.5 rise=3 targ v(Dout_0) val=0.5 cross=1 td=4n
* Power metric: average supply current over full run
meas tran iavg_vdd avg i(Vdd) from=0n to=12n
let pavg_mw = -iavg_vdd * 1.0 * 1e3
let hex_5ns = D3_at_5ns*8 + D2_at_5ns*4 + D1_at_5ns*2 + D0_at_5ns
let hex_7ns = D3_at_7ns*8 + D2_at_7ns*4 + D1_at_7ns*2 + D0_at_7ns
let f_eq_ghz = 1e-9 / t_clk_to_dout
print t_clk_to_dout iavg_vdd pavg_mw f_eq_ghz hex_5ns hex_7ns
.endc
```

8.2 Top-level functional and timing result

The full validation terminal output is shown as Fig. 37; the corresponding multi-signal SPICE transient is shown as Fig. 42 in §, where it sits next to the timing analysis it quantifies.

```

=====
                                16x4 SRAM FUNCTIONAL VALIDATION
=====
Write phase: 0-4 ns
Read phase : 4-8 ns
Extra write: 8-10 ns

d0_at_5ns      = 9.999427e-01
d1_at_5ns      = 4.868107e-04
d2_at_5ns      = 9.999427e-01
d3_at_5ns      = 4.868107e-04
d0_at_7ns      = 4.868106e-04
d1_at_7ns      = 9.999427e-01
d2_at_7ns      = 4.868106e-04
d3_at_7ns      = 9.999427e-01
t_clk_to_dout  = 1.106501e-10 targ= 4.215650e-09 trig= 4.105000e-09
iavg_vdd       = -2.136728e-05 from= 0.000000e+00 to= 1.200000e-08
=====
Readback @5 ns (expect 0x5 = 0101)
=====
d3_at_5ns = 4.868107e-04
d2_at_5ns = 9.999427e-01
d1_at_5ns = 4.868107e-04
d0_at_5ns = 9.999427e-01
hex_5ns = 5.004582e+00
=====
Readback @7 ns (expect 0xA = 1010)
=====
d3_at_7ns = 9.999427e-01
d2_at_7ns = 4.868106e-04
d1_at_7ns = 9.999427e-01
d0_at_7ns = 4.868106e-04
hex_7ns = 1.000186e+01
=====
Performance Metrics
=====
t_clk_to_dout = 1.106501e-10
iavg_vdd = -2.13673e-05
pavg_mw = 2.136728e-02
f_eq_ghz = 9.037497e+00
=====
Validation Status
=====
Address 0 readback: PASS
Address 1 readback: PASS
Timing measurement: VALID
=====

```

Figure 37: Top-level extracted-netlist validation terminal output (*media/terminal/final-terminal-out.png*, generated by the *.control* block above). Both readbacks PASS, $t_{\text{CLK} \rightarrow D_{\text{out}}} = 1.1065 \times 10^{-10} \text{ s}$, $I_{DD,\text{avg}} = -2.1367 \times 10^{-5} \text{ A}$ ($P = 21.367 \mu\text{W}$), and the equivalent frequency derived from the $\text{CLK-to-}D_{\text{out}}$ delay is $f_{\text{eq}} = 9.037 \text{ GHz}$, $\sim 18.07 \times$ the 500 MHz spec floor.

Readback @ 5 ns : 0x5 (PASS)	Readback @ 7 ns : 0xA (PASS)
$t_{\text{CLK} \rightarrow D_{\text{out}}} : 110.65 \text{ ps}$	$P_{\text{avg}} : 21.367 \text{ } \mu\text{W}$
Sustainable max CLK f_{max} (W/W/R/R sweep) : 4.571 GHz @ 218.75 ps	
$f_{\text{eq}} = 1/t_{\text{CLK} \rightarrow D_{\text{out}}} \text{ (single-edge latency)} : 9.037 \text{ GHz}$	

8.3 Sustainable max CLK frequency sweep (find_fmax.py)

The validation `.control` block of § reports the *single-edge* CLK-to- D_{out} access time (110.65 ps) and its inverse $f_{\text{eq}} = 1/t_{\text{CLK} \rightarrow D_{\text{out}}} \approx 9.037 \text{ GHz}$. That is a clean access-time figure but it is *not* the same quantity as a sustainable back-to-back operating frequency: at very short CLK periods the precharge / non-overlap / write-driver settling windows have not yet finished by the time the next CLK edge arrives, even though a single read edge resolves D_{out} in $\sim 110 \text{ ps}$.

To measure the spec-defined “max CLK frequency (i.e. min CLK period)”, the helper script `spice/find_fmax.py` was run against the same extracted top deck. The script binary-searches the CLK period: at each candidate T it scales every PWL source (CLK, $A[3:0]$, $D_{\text{in}}[3:0]$, WE) to a 4-cycle write/write/read/read schedule, reruns `ngspice -b`, samples $D_{\text{out}}[3:0]$ at the two read instants, thresholds at 0.5 V, and accepts the period only if both readbacks (0x5 and 0xA) succeed. The search bracket was [0.10, 2.00] ns with a 5 ps tolerance. For repeatable command-line runs with a formatted banner, optional spinner, and automatic JSON capture, the repository includes the wrapper script `scripts/run-find-fmax.sh`, which calls `find_fmax.py` with the same arguments used to produce the reported evidence. Figure 38 summarizes the search loop; Figure 39 shows a saved `ngspice` transcript for the same flow.

Regression and screening rigor. Beyond the headline bisection, the repository includes a parallel driver (`spice/parallel_cell_mix.py`) that composes dozens of *integer* $L = 1 / L = 2$ width recipes (per-library cell scales), runs `find_fmax.py` concurrently, and logs CSV outcomes. Recent sweeps varied clock-local devices (`nor2`, `inv_X8`, `inv`), replica and decode stacks, precharge, flip-flop buffers, and dummy-column loads. On the current `top.spi` export, every screened point returned the same sustainable closure $T_{\text{min}} = 0.21875 \text{ ns}$ ($f_{\text{max}} = 4.571 \text{ GHz}$), which is strong evidence that the limiter is a shared timing envelope (precharge / non-overlap / bitline budget) rather than any single tunable cell instance. The flowchart in Fig. 38 is the canonical depiction of the acceptance test inside each of those jobs.




Figure 38: *find_fmax.py*: bisection on CLK period with functional readback as the acceptance test (§).

```

=====
Full SRAM W/W/R/R Sustained f_max Sweep Summary
=====
Deck                : spice/top.spi
Sustained f_max     : 4.571429 GHz
T_min               : 218.7500 ps (0.218750 ns)
Margin vs 500 MHz spec : 9.1429x
Steady-state verify  : PASS (32 macros, 128 CLK cycles)
t_CLK→Dout (single-edge): 101.5406 ps
f_eq = 1/t_CLK→Dout   : 9.848277 GHz
P_avg (from i(Vdd))   : 149.7311 μW @ VDD=1.0 V
P_avg window         : 0 → 0.984375 ns (sweep deck)
FOM (access delay, sweep): 7.4102e-22 [Area=8]
FOM (cycle @ T_min, sweep): 3.4391e-21 [Area=8]
Pattern              : W/W/R/R addr0=0x5 addr1=0xA functional readback @0.5V

End:    2026-05-09 02:17:25
Elapsed: 01:39 (mm:ss)
Saved JSON: spice/find_fmax_last.json

```

Figure 39: `find_fmax.py` closure on the shipped `top.spi`: best passing period 0.21875 ns ($f_{\max} = 4.571$ GHz), 9.14× vs. 500 MHz spec, $t_{\text{CLK} \rightarrow D_{\text{out}}} \approx 102$ ps at $T_{\min}$, bisection tolerance ± 5 ps, and steady-state verification PASS (32 W/W/R/R macros, 128 CLK cycles). File: `media/terminal/fmax.png` (refresh after re-running the script for a matching screenshot).

Sustainable max CLK		$T_{\min} : 218.75$ ps	$f_{\max} : 4.571$ GHz
both readbacks: PASS at $T_{\min}$		margin vs. 500 MHz spec: 9.14×	

The $T_{\min} \approx 218.75$ ps result is bounded by the precharge + non-overlap + write settling envelope, not by the raw read access path alone; the single-read $t_{\text{CLK} \rightarrow D_{\text{out}}}$ at the bisected fast period is on the order of 100–110 ps, while the 500 MHz validation deck (§) reports 110.65 ps on the address-0 read edge. The two frequency claims therefore describe different things and are both kept in the report: $f_{\text{eq}} \approx 9.037$ GHz is the inverse of the read-access latency that feeds the FOM *delay* term (§), while $f_{\max} = 4.571$ GHz is the sustainable cycle frequency that the spec calls out as the “max CLK frequency”.

9 Critical-Path Timing Analysis

9.1 Read path (worst case)

The read critical path is from CLK rising edge to D_{out} valid. Components (stacked proportionally in Figure 40):

Stage	Delay (ps)
CLK $\rightarrow$ Φ_2 rise (non-overlap generator)	40
Address register CLK-to-Q	25
Predecoder NAND2 + INV	50
Row NAND3 + WL driver INV	50
Wordline rise (3 fF load, sized driver)	20
BL discharge to 100 mV (post-extraction $C_{BL} = 13$ fF)	37
SAE generation (replica trip + tunable delay)	15
StrongARM resolve to full rail	45
Output latch + buffer	30
Read total	312

Table 3: Read-path delay breakdown, post-extraction.




Figure 40: Read-path delay stack (ps per segment; total 312ps analytic budget, Tab. above). Segment widths are proportional to the table.

The replica column SAE generation overlaps the BL discharge: SAE fires when the dummy BL has discharged by ~ 100 mV, which is roughly contemporaneous with the real BL crossing the SA-required differential. The 15 ps line item is the additional tunable delay between the dummy-BL trip and SAE assertion, not the dummy-BL discharge itself (which is already counted in the BL-discharge line).

9.2 Write path (worst case)

Figure 41 stacks the same stages proportionally to their ps contributions.

Stage	Delay (ps)
CLK $\rightarrow$ Φ_2 rise (non-overlap generator)	40
Address register CLK-to-Q	25
Predecoder + row NAND + WL driver	110
Wordline rise	20
BL/BLB driven to rail by write driver	55
Cell storage node flip	80
Write total	330

Table 4: Write-path delay breakdown. The BL slew and the cell flip are partially in series because the cell cannot flip until BL/BLB have crossed the inverter trip point of the cross-coupled inverters.




Figure 41: Write-path delay stack (330 ps total; segment legend matches Tab. above).

The two paths overlap heavily in their early stages; the difference is in the back end. Read is gated by SA decision time (~ 45 ps after BL differential develops); write is gated by the cell-flip time (~ 80 ps after BL/BLB develop). Write is slightly slower in this design because we are pushing C_{BL} to full rails rather than just to $\Delta V = 100$ mV.

9.3 Access delay and operating point

The stage-by-stage analysis above provides a conservative bound on the worst-case access delay. To report the system-level timing in the same form used by the final top-level validation, we additionally measure the direct CLK -to- D_{out} delay in the full extracted SRAM deck (§). In the final top-level validation run, this measured delay is $t_{CLK \rightarrow D_{out}} = 110.65$ ps.

For a safe operating point, the CLK period must accommodate the active phase of the operation, the precharge/recovery phase, and non-overlap guard time:

$$T_{\min} = T_{\Phi_2, \text{active}} + T_{\Phi_1, \text{precharge}} + 2 \cdot t_{\text{nonoverlap}}.$$

The cycle-stepped sweep in § closes this budget empirically and confirms the sustainable steady-state operating point at $T_{\min} \approx 218.75$ ps ($f_{\max} \approx 4.571$ GHz). At the same fast period the single-edge access $t_{CLK \rightarrow D_{out}}$ remains on the order of 100–110 ps, while the 500 MHz validation deck reports 110.65 ps, so the ratio $T_{\min}/t_{CLK \rightarrow D_{out}} \approx 2.0$ is dominated by the precharge + non-overlap envelope rather than the read access path alone. This report uses the top-level measured access delay (110.65 ps) as the delay term for the final Figure of Merit (§) to keep the reported FOM tied directly to observed read-edge behavior; an alternative cycle-time-based FOM is noted inline in §.




Figure 42: Top-level SRAM transient (*media/waveforms/sram_full_validation.png*) generated by the `.control` block of §. Traces, from top: CLK, WE, A₁, A₀, D_{in,0}, D_{out,0}, D_{out,1}, D_{out,2}, D_{out,3} (offset for visibility). The two read windows at ~ 5 ns and ~ 7 ns produce D_{out} = 0x5 and D_{out} = 0xA respectively, matching the readbacks in Fig. 37; the CLK-to-D_{out} delay measured on the address-0 read edge is 110.65 ps.

10 Power Measurement

10.1 Methodology

The average power is measured directly in the top-level extracted deck by integrating the supply current and dividing by the measurement window. Power is then computed as $P = V_{DD} \cdot I_{DD,avg}$. The final top-level validation run (§) reports $I_{DD,avg} = 21.367 \mu\text{A}$ at $V_{DD} = 1$ V, corresponding to $P_{avg} = 21.367 \mu\text{W}$ (boxed below). The average is taken over the full 12 ns tran window in `top.spi`, which includes idle phases between stimulus edges, so it is a time-mean current rather than a peak switching current.

$I_{DD,avg}$ (12 ns window) : 21.367 μ A	V_{DD} : 1.000 V
$P_{avg} = V_{DD} \cdot I_{DD}$: 21.367 μ W	

10.2 Power decomposition (qualitative)

The average power has four dominant components per cycle; Figure 43 shows the same split as a stacked fJ bar.

Component	Energy / cycle
BL/BLB precharge recovery (4 columns, $\sim$ half swing each)	~ 16 fJ
Write driver active swing	~ 7 fJ
Cell storage-node flip energy (~ 4 cells flip per cycle in stress test)	~ 10 fJ
Decoder + clock + WL drivers + replica column + sized output buffers	~ 10 fJ
Total per cycle (at $f_{CLK} = 500$ MHz)	~ 42.7 fJ

Table 5: Energy breakdown per cycle during the write-stress test (scaled so the totals match the measured $E_{cyc} = P_{avg}/f_{CLK} = 21.367 \mu\text{W}/500 \text{ MHz} \approx 42.7 \text{ fJ}$). Component splits are qualitative; the reported power is the simulated $V_{DD} \cdot I_{DD,avg}$ result.




Figure 43: Qualitative per-cycle energy (fJ per segment; legend matches Tab. above).

This decomposition is intended to build intuition about where the energy is spent in a cycle. The authoritative reported power for the completed design is the measured $V_{DD} \cdot I_{DD,avg}$ value from the top-level validation run (Fig. 37).

11 Figure of Merit

11.1 Numerical FOM

Inputs to the FOM, from the preceding sections:

- BitcellArea = 8 (in units of W_{min}).
- Power = 21.367 μ W = 2.1367×10^{-5} W (top-level measured).
- Delay = 110.65 ps (top-level measured); Delay² = 1.224×10^{-20} s².

$$\text{FOM} = 60 \cdot 8 \cdot 2.1367 \times 10^{-5} \cdot \left(1.1065 \times 10^{-10}\right)^2 \approx 1.256 \times 10^{-22}.$$

Access-time vs. cycle-time FOM. The number above uses the single-edge read access $t_{\text{CLK} \rightarrow D_{\text{out}}} = 110.65$ ps as the delay term, which is the standard SRAM convention and matches the “read access time” the spec asks for. For completeness, an alternative cycle-time-based FOM that uses the sustainable minimum period $T_{\text{min}} = 218.75$ ps from § as the delay term gives $60 \cdot 8 \cdot 2.1367 \times 10^{-5} \cdot (2.1875 \times 10^{-10})^2 \approx 4.91 \times 10^{-22}$. Both formulations meet the spec floor by large margins; this report carries the access-time value forward as the headline FOM.

What the scalar FOM is (and is not). The project formula multiplies declared inputs: bitcell width sum, measured average supply power, and a delay. Here P_{avg} is the time-mean $V_{DD} \cdot I_{DD}$ over the 12 ns top-level validation transient at $f_{\text{CLK}} = 500$ MHz (§), while D is the single-edge CLK-to- D_{out} latency from the same deck (§). Sustainable max CLK $f_{\text{max}} \approx 4.571$ GHz is established separately by the `find_fmax.py` sweep (§) and is not used as the delay factor in the headline FOM. The headline value should therefore be read as the rubric composite from those cited measurements—not as a literal single-point energy–delay product at sustained f_{max} , and not as average power at maximum switching throughput (which would require a dedicated long transient or average current measurement matched to that activity). The alternate cycle-time FOM above likewise keeps the same P_{avg} for comparison only; remeasuring P under a T_{min} macro would be needed for a fully self-consistent throughput-matched composite.

11.2 Sensitivity table

Figure 44 plots relative FOM for each quoted alternative; the table keeps delay and power columns. The following table shows how the FOM would change under specific alternative design choices, holding other parameters fixed at their chosen-design values. This documents the value of each architectural decision.

Variant	W	Delay (ps)	P (μW)	FOM (rel.)
This design (6T 2-1-1, no mux, StrongARM, replica)	8	111	21.367	1.00
8T cell (with read port)	13	300	130	3.75
6T 1-1-1 (CR=PR=1)	6	—	—	<i>fails read SNM</i>
Static differential SA	8	430	240	8.75
Column-muxed 4×16	8	380	130	3.70
Hardcoded SAE delay (no replica)	8	360	122	3.12

Table 6: FOM sensitivity to design choices. Column W : bitcell summed width. Values for alternatives are estimated from the same circuit-level analysis used for the chosen design and are intended as relative comparisons, not high-precision predictions.




Figure 44: FOM sensitivity (Tab. above): horizontal extent is proportional to relative FOM; static SA dominates.

The largest single penalty is the static sense amplifier: $8.7\times$ the chosen-design FOM, driven by both higher delay and higher power. Column muxing and an 8T cell each cost $\sim 3.7\times$. The replica column buys $\sim 3.1\times$ over a hardcoded SAE chain. (The chosen-design baseline is now small enough in delay² that any alternative carrying 60–90% more delay is penalized superlinearly.)

12 Validation Methodology and Summary

12.1 Test methodology

All validation tests are run on the full extracted netlist. Functional correctness is demonstrated using a conservative 500 MHz clock ($T_{CLK} = 2\text{ ns}$) with PWL-driven inputs for CLK, $A[3:0]$, $D_{in}[3:0]$, and WE . Expected D_{out} is sampled at fixed read instants (5 ns and 7 ns) corresponding to the two read phases of the validation sequence. Two distinct frequency claims are reported and should not be conflated: the *single-edge access* latency $t_{CLK \rightarrow D_{out}} = 110.65\text{ ps}$ from the validation `.control` block (whose inverse is the $f_{eq} \approx 9.037\text{ GHz}$ figure that feeds the FOM delay term), and the *sustainable max CLK frequency* $f_{max} \approx 4.571\text{ GHz}$ (period $T_{min} = 218.75\text{ ps}$) measured by the cycle-stepped sweep in § via `spice/find_fmax.py` (or the convenience harness `scripts/run-find-fmax.sh`). The latter is the spec’s “max CLK frequency / min CLK period” quantity. All in-line evidence (SPICE transients and NGSPICE `.measure` terminal output) for each block has already been placed alongside the corresponding claim in §–§; this section consolidates the test-case set itself plus the cross-reference summary table. Parallel integer-width sweeps (`parallel_cell_mix.py`, logged CSVs under `spice/`) were additionally used as a regression screen; on the current extracted deck they reproduce the same T_{min} closure reported above.

12.2 Test cases

1. **Write-then-read at every address, both data patterns.** For each of the 16 addresses, write 0000 then read; check $D_{out} = 0000$. Repeat for 1111. Total: 32 cycles per pattern. Validates basic read and write functionality.
2. **Walking ones.** Write 0001, 0010, 0100, 1000 to consecutive addresses, then read each back. Repeat for walking zeros (1110, 1101, 1011, 0111). Validates that columns are independent and BL/BLB are not coupled.
3. **Read-disturb stress.** Write 0101 to address 0. Read address 0 1000 times consecutively. Verify each read returns 0101. Validates the cell ratio choice (this is the test that fails for $CR = 1$).
4. **Half-select check.** Write 0000 to address 5. Verify the contents of addresses 4, 6, 7, and 13 (each shares some address bits with 5 through the predecoder, so any half-select bug would manifest here). Validates that unselected rows are isolated.
5. **Back-to-back write/read same address.** Write 1010 to address 7; on the very next cycle, read address 7. Verify $D_{out} = 1010$. Validates that BL recovery during Φ_1 completes in time.
6. **Mid-cycle input change.** During Φ_2 of a write cycle, flip D_{in} from 1100 to 0011. Verify the edge-triggered value (1100) is what was written. Validates the input registers.
7. **Sustainable- f_{max} sweep (W/W/R/R).** The spec-defined max clock frequency is measured using the cycle-scaled W/W/R/R macro sweep in §. The script `spice/find_fmax.py` binary-searches the minimum CLK period, scaling *all* PWL sources each trial, and accepts a candidate

period only if both readbacks (0x5 and 0xA) pass at a 0.5 V threshold. The optional post-search steady verify runs many consecutive macros at $T_{\min}$ to strengthen the sustained claim.

12.3 Cross-reference of measured evidence

The following table consolidates every measured value cited in the report, in the order it is introduced, and points each entry at the exact in-repo terminal/waveform PNG that was used to produce it.

Test / metric	Measured value	Evidence (in-repo path)
DFF CLK-to- Q rise / fall / avg	42.0 / 91.9 / 67.0 ps	media/terminal/dffvalidation.png
Bitcell Test 1: $t_{\text{write},1 \rightarrow 0}$	9.87 ps	media/terminal/test1-sramwrite1-0.png
Bitcell Test 2: $t_{\text{write},0 \rightarrow 1}$	21.8 ps	media/terminal/test2-sram-write0->1.png
Bitcell Test 3: max V_Q during read (store 0)	95.4 mV	media/terminal/test3-read-stability.png
Bitcell Test 4: max V_Q during read (store 1)	95.4 mV	media/terminal/test4-read-stability.png
Bitcell Test 5: min V_Q over 5 ns retention	0.9999905 V	media/terminal/test5-data-retention sram.png
Precharge: $t_{\text{pch,BL}}$ (0→0.95 V)	179.8 ps	media/terminal/prechargeeq results.png
Precharge: equalize time t_{equalize}	158.8 ps	media/terminal/prechargeeq results.png
Write driver (13 fF load): $t_{\text{write},\text{rise}}$	167.3 ps	media/terminal/write-driver-load.png
Write driver (13 fF load): $t_{\text{write},\text{fall}}$	142.9 ps	media/terminal/write-driver-load.png
Top-level $t_{\text{CLK} \rightarrow D_{\text{out}}}$ (single-edge access)	110.65 ps	media/terminal/final-terminal-out.png
Top-level $I_{DD,\text{avg}}$ / P_{avg}	21.367 μA / 21.367 μW	media/terminal/final-terminal-out.png
Top-level readback @ 5 ns / @ 7 ns	0x5 / 0xA (PASS)	media/terminal/final-terminal-out.png
Delay-derived equivalent frequency $f_{\text{eq}} = 1/t_{\text{CLK} \rightarrow D_{\text{out}}}$	9.037 GHz @ 110.65 ps	media/terminal/final-terminal-out.png
Sustainable max CLK frequency f_{max} (W/W/R/R sweep)	4.571 GHz @ 218.75 ps	media/terminal/fmax.png, §

Table 7: Cross-reference of measured values used in the report. Values are taken from NGSPICE .measure terminal captures or the cycle-scaled sweep script; where reported values are rounded for readability, the underlying source values are preserved in the linked evidence files. The in-repo paths can be opened verbatim against the submission tree.

12.4 Spec compliance recap

Every requirement listed in proj2.pdf is met by margin:

- **16×4 SRAM, 4-bit address, 4-bit $D_{\text{in}}/D_{\text{out}}$, WE, single CLK** — top-level netlist (§).
- $V_{DD} \leq 1 \text{ V}$ — $V_{DD} = 1.000 \text{ V}$ (vdd vdd 0 DC 1.0).
- **Synchronous**, $f_{\text{max}} \geq 500 \text{ MHz}$ — sustainable max CLK frequency from the cycle-stepped find_fmax.py sweep $f_{\text{max}} \approx 4.571 \text{ GHz}$ (min period $T_{\min} \approx 218.75 \text{ ps}$), $\sim 9.14\times$ over the spec floor; the delay-derived $f_{\text{eq}} = 1/t_{\text{CLK} \rightarrow D_{\text{out}}} \approx 9.037 \text{ GHz}$ is also reported for the single-edge read access (§).
- **Read or write completes within one CLK period** — $t_{\text{CLK} \rightarrow D_{\text{out}}} = 110.65 \text{ ps}$ single-edge access; cell-level write times 9.87/21.8 ps; spec is met at any cycle $T \geq T_{\min} = 218.75 \text{ ps}$.
- **Bitcell area metric** — $8 W_{\min}$ summed transistor width per cell (§).

- **Power measurement methodology** — average V_{DD} current integrated over the validation transient (§); reported $P_{\text{avg}} = 21.367 \mu\text{W}$.
- **FOM reported** — $\text{FOM} = 60 \cdot 8 \cdot P \cdot D^2 \approx 1.256 \times 10^{-22}$ (§).

13 Summary of Design Metrics

Metric	Value
Process	PTM 22 nm High-Performance
Supply voltage	1.0 V
Capacity	16 words $\times$ 4 bits
Bitcell topology	6T
Bitcell device count	6
Cell ratio (CR)	2
Pull-up ratio (PR)	1
Bitcell summed transistor width	$8 W_{\min}$
Array organization	16 rows $\times$ 4 columns; no column mux
Sense amplifier	Clocked StrongARM regenerative latch
SAE generation	Replica/dummy bitline column (self-timed)
Clocking scheme	Two-phase non-overlapping from single CLK
Worst-case access delay (stage-by-stage, post-extr.)	330 ps (write-limited)
Minimum CLK period (with margin)	219 ps
Measured CLK-to- D_{out} (single-edge access)	110.65 ps
Sustainable max CLK frequency $f_{\max}$ (W/W/R/R sweep)	4.571 GHz @ 218.75 ps
Equivalent frequency $f_{eq} = 1/t_{\text{CLK} \rightarrow D_{\text{out}}}$ (single-edge)	9.037 GHz
Specification frequency requirement	≥ 500 MHz (sustainable: $9.14\times$ margin)
Average power, top-level run	21.367 μW
DFF CLK-to- Q average (t_{cq})	67.0 ps
Bitcell write 1 $\rightarrow$ 0 / 0 $\rightarrow$ 1 time	9.87 ps / 21.8 ps
Read disturb max V_Q (Tests 3, 4)	95.4 mV (vs. 0.4 V threshold)
Retention drop over 5 ns	$< 10 \mu\text{V}$
Functional readback @ 5 ns / @ 7 ns	0x5 / 0xA (PASS)
FOM = $60 \cdot A \cdot P \cdot D^2$	$\approx 1.256 \times 10^{-22}$

Table 8: Final summary of all design metrics.

Appendix A — Repository Map

The full source tree (SPICE deck, Electric library, plotting scripts, and the LaTeX manuscript) is mirrored at github.com/tmarhguy/64b-sram. The complete simulation deck and Electric library backing every figure and number in this report are organized as follows in the submission tree:

- `spice/top.spi` — top-level netlist + validation `.control` block (excerpted in §).
- `spice/find_fmax.py` — cycle-scaled $f_{\max}$ sweep driver. Binary-searches the CLK period under a 4-cycle write/write/read/read schedule against the live extracted `top.spi`; closes at $T_{\min} = 218.75$ ps ($f_{\max} = 4.571$ GHz), reproducing the sustainable max-CLK number cited in §.
- `spice/sram_6t_cell.spi`, `spice/dummy_bitcell.spi`, `spice/row_4.spi`, `spice/array_16x4.spi`, `spice/replica_column.spi` — core array.
- `spice/precharge_eq.spi`, `spice/write_driver.spi`, `spice/sense_amp.spi`, `spice/output_latch.spi`, `spice/column_slice.spi` — column periphery.
- `spice/predecoder_2to4.spi`, `spice/predecoder_4to16.spi`, `spice/nonoverlap_clk_gen.spi`, `spice/dff.spi` — control.
- `spice/proj2.jelib` — Electric library producing every schematic in `media/`.
- `media/waveforms/*.png` — NGSPICE transient exports cited inline (matching SVG masters under `spice/waveforms/`).
- `media/terminal/*.png` — saved NGSPICE `.measure` terminal captures cited inline as evidence.

Appendix B — Electric Hierarchy Checklist

For layout and verification cross-reference:

- Core: `sram_6t_cell`, `dummy_bitcell`, `row_4`, `array_16x4`, `replica_column`.
- Decoder / clock: `predecoder_2to4`, `predecoder_4to16`, `nonoverlap_clk_gen` (clock-gen-2phase).
- Column periphery: `precharge_eq`, `write_driver` (incl. `tristate_inv`), `sense_amp` (StrongARM), `output_latch`, `column_slice`.
- Primitives: `inv`, `inv_X4`, `inv_X5`, `inv_X8`, `invX2`, `invX2_2`, `nand2`, `nand3`, `nor2`, `tgate`, `tgate2`, `dff`.
- Top-level: `input registers (9×dff) + predecoder_4to16 + clock-gen-2phase + 4×column_slice + replica_column + nand2+inv_X4 (for WD_EN) = top`.

Academic Integrity Statement

I, **Tyrone Marhguy**, certify that I have complied with the University of Pennsylvania’s Code of Academic Integrity in completing this project.